

Python Web Development

By Dr. Tanima

Quick URLs:

- <http://python.org/>
- <http://webware.sourceforge.net/>
- <http://www.python.org/cgi-bin/moinmoin/WebProgramming>
- <http://www.djangoproject.com>
- <https://pypi.python.org>
- <https://www.djangopackages.com>

What is persistence?


- Data lives longer than programs
 - Files, pipes, relational databases, etc.
 - But, discontinuity of representation : object serialization
- Orthogonal persistence
 - Automatic management of program state
 - Independent of data type or longevity
 - Allow programmer to focus on application data model

Python approach

- Goals
 - Minimize changes to existing programs
 - Work with standard interpreter
- Zope Object Database (ZODB)
 - Support Zope application server
 - Originally targeted for web development
 - Zope (Z Object Publishing Environment)
 - Any object reachable from ZODB root is persistent
 - Other App Server is Webware
- Separate database and storage layers
 - Store object's data in formatted text file, such as a CSV file.
 - Or can use a relational database, such as Gadfly, MySQL, PostgreSQL, or DB2.
 - These file formats and databases are well established, and Python has robust interfaces for all of these storage mechanisms.

Persistence by reachability

- Any object reachable from ZODB root is persistent



ZOPE: A simple example

```
from Persistence import Persistent
from Transaction import get_transaction
from ZODB.FileStorage import DB
```

```
class Counter(Persistent):
    _value = 0
    def inc(self):
        self._value += 1
```

```
def main():
    fs = DB("data.fs")
    conn = db.open(); root = conn.root()
    obj = root["myobj"] = Counter()
    get_transaction().commit()
    obj.inc()
    get_transaction().commit()
```


Object serialization

- Standard pickle library
 - Serializes arbitrary object graph
 - Raises `TypeError` for sockets, files, etc.
 - Instance vars serialized via dictionary
- Hooks to define custom state
 - `__getstate__()` / `__setstate__()`
 - Persistent mixin ignores `_v_` attributes

Pickling persistent objects

- Stores objects in separate records
 - Persistent objects pickled as oid + class
 - Works with cache to maintain identity
- Handling non-persistent objects
 - Copied into record of containing object
 - Sharing by persistent objects is problematic

Object identity / caching

- Cache maintains oid  object mapping
 - Guarantees only one copy of object
 - Unpickler loads all referenced objects
- Ghost objects
 - A webkit web client written in python
 - Only Persistent header initialized
 - No instance state loaded
 - State loaded on first object access
 - `from ghost import Ghost`
- LRU cache of recent objects

Attribute access handlers

- Persistent implements C wrappers
 - Override `tp_getattro`, `tp_setattro` slots
 - Mediate access to instance variables
 - Crucial Python feature

Transactions

- Supports multiple threads, processes
 - Independent database connections
 - Updates visible at transaction boundaries
- Optimistic concurrency control
 - When conflict occurs, abort and retry
- On error, abort to restore consistency
 - Reverts to last saved state

Concurrency and conflicts

- Invalidations sent at commit time
 - Clients process at transaction boundaries
- Conflicting transactions aborted
 - Write conflict at commit time
 - Read conflict on object access
- Application must retry on conflict
 - Can use generic wrapper
 - Can define conflict resolution method

Other features

- Undo support
 - Storage stores multiple revisions
 - Transactional undo reverts to earlier state
- BTrees: efficient persistent containers
- Storing code in database

Limitations

- Schema evolution
 - Must code manually in `__setstate__()`
- Database management
 - Manual `pack()` to remove revisions, do GC
- Sharing of non-persistent objects
- Integration with legacy code
 - Multiple inheritance helps
 - Factory classes